

MASTER'S DEGREE IN PHYSICS
Academic Year 2020-2021
QUANTUM INFORMATION

Student: Giorgio Palermo
Student ID: 1238258
Date: November 2, 2020

EXERCISE 4

In this report I will describe the details on how I tested the performance of my machine relative to matrix multiplication using Fortran subroutines and Python scripts; I will briefly evaluate trend of the data I obtained for matrices of sizes spanning from 10 to 2500 units.

1 Theory

The basic concepts I used to build this program are relative to Fortran functions and subroutines, to the Fortran input-output procedures relative to text files and to the communication between Fortran programs and bash. I also used some basic commands of Python scripting language, including functions from NumPy and Os modules.

Code Development

To develop the code I used to solve this exercise I started from the `MatTest` subroutine I implemented for `ex2` and I modified it. In the first place `MatTest` performed a matrix multiplication test in an almost automatic mode: given the maximum size of the matrix it tested the computation time for different algorithms increasing the matrix size by 100 units steps. For this exercise the test couldn't be automatic anymore: the matrix size had to be specified arbitrarily iteration by iteration. For this reason I implemented a new version of `MatTest` (reported below) with the following characteristics:

- it measures the computation time for a $N \times N$ matrix multiplication between random matrices, with N given as input
- it outputs the results on different files for each method; the file name is passed as input; an optional live output of the results is implemented
- the result is passed to the calling program using a one-dimensional array containing matrix size and time results for each iteration.

```
1 subroutine MatTest(filename, size_input, verbose, result)
2     use Functions
3     use Debug
4     implicit none
5     ! Local scalars
6     character(*), intent(in) :: filename
7     character(*), intent(in) :: verbose
8     character(len=50) :: filename_in, filename_alpha, filename_bravo, message
9     double precision, intent(in) :: size_input
10    integer :: size
11    integer :: nn=5, ii=0, status
12    integer :: oo=6 ! to suppress screen output set oo=something
13    double precision :: start=0, finish=0, sum=0
14    ! Local Arrays
15    double precision, dimension(4) :: result
16    double precision, dimension(:, ::), allocatable :: A,B,C,C1,C2
17    double precision, dimension(1,3) :: time
18    character(len=30), dimension(:, ::), allocatable :: args
19    size=floor(size_input)
20    ! Select if verbose
21    select case(verbose)
22    case("y")
23        oo=6
24    case("n")
```

```

25         oo=3456
26     case default
27         oo=3456
28     end select
29
30     write(oo,*)
31     write(oo,*) "    *** Matrix multiplication test program ***  "
32
33     ! filename_in = filename
34     filename_in = filename // "_ByRows" // ".txt"
35     filename_alpha = filename // "_ByCols" // ".txt"
36     filename_bravo = filename // "_Intrinsic" // ".txt"
37
38     ! Testing section
39     open(unit=10,file=filename_in,action='write',position="append",status='unknown',
,iostat=status)
40     open(unit=20,file=filename_alpha,action='write',position="append",status='
unknown',iostat=status)
41     open(unit=30,file=filename_bravo,action='write',position="append",status='
unknown',iostat=status)
42     write(oo,*)
43     write(oo,*) "Running test..." ! courtesy message
44     write(oo,*) "Size      ", "      ByRows[s]  ", "      ByCols[s]    ", "      Intrinsic[s]
"
45
46     allocate(A(size,size),B(size,size),C(size,size),C1(size,size),C2(size,size))
47     call random_number(A)
48     call random_number(B)
49
50     call cpu_time(start)
51     call LoopMult(A,B,C)      ! by rows
52     call cpu_time(finish)
53     time(1,1)=finish-start
54
55     call cpu_time(start)
56     call LoopMultColumns(A,B,C1) ! by columns
57     call cpu_time(finish)
58     time(1,2)=finish-start
59
60     call cpu_time(start)
61     call IntrinsicMult(A,B,C2) ! Intrinsic
62     call cpu_time(finish)
63     time(1,3)=finish-start
64     deallocate(A,B,C,C1,C2)
65
66     write(oo, '(i5,3G15.5)') size,time(1,1),time(1,2),time(1,3)
67     write(10, '(i10, G15.5)') size, time(1,1) ! Write output on different files
68     write(20, '(i10, G15.5)') size, time(1,2)
69     write(30, '(i10, G15.5)') size, time(1,3)
70     result(1) = dble(size)
71     result(2) = time(1,1)
72     result(3) = time(1,2)
73     result(4) = time(1,3)
74     write(oo,*) "Done"
75     close(10,iostat=status)
76     close(20,iostat=status)
77     close(30,iostat=status)
78
79     end subroutine MatTest

```

To import the grid of points containing the matrix sizes to test I implemented the `ReadGrid` subroutine: this piece of code reads as inputs the name of the file to read from (`filename`) and the grid size (`grid_dim`); the output is a one-dimensional array containing the points to test:

```

1  subroutine ReadGrid(filename,grid_dim, grid)
2      implicit none
3      character(*), intent(in) :: filename
4      character(len=100) :: msg
5      integer, intent(in) :: grid_dim
6      double precision, dimension(:), allocatable, intent(out) :: grid
7      integer :: ii=1, ios
8
9      open(unit=100,file=filename,iostat=ios,iomsg=msg)
10     if(ios/=0) then

```

```

11         write(*,*) msg
12         stop
13     end if
14
15     allocate(grid(grid_dim))
16     read(100,*,iostat=ios, end=997,iomsg=msg) grid
17     997 print*, "Grid read with exit status: ", ios
18     ! print*, grid(:)
19     close(100)
20 end subroutine ReadGrid

```

The real testing is implemented in `program cos`. This program calls repeatedly the `MatTest` subroutine, giving as input a different value of N for each call. The grid is read from file using the `ReadGrid` function, while the output results are collected in a two dimensional array `results_all`. In order to be easily scripted with python, I wrote this program to take an input from bash (`filename`), which I used to pass to the program the name of the file to store the results in.

```

1 program cos
2     use Functions
3     use cosmod
4     implicit none
5
6     double precision, dimension(:,,:), allocatable :: results_all
7     double precision, dimension(:), allocatable :: grid
8     double precision, dimension(4) :: result
9     character(len=30), dimension(:), allocatable :: args
10    integer grid_dim,ii,ios, num_args
11    character(len=30) :: filename
12
13    num_args = command_argument_count()
14    allocate(args(num_args))
15    call get_command_argument(1,args(1))
16    print*, args(1)
17    filename = "result"// trim(args(1))// ".dat"
18    print*, filename
19    grid_dim=10
20    allocate(grid(grid_dim))
21    allocate(results_all(grid_dim,4))
22
23    call ReadGrid("grid.dat", grid_dim,grid)
24    open(unit=78,file=filename,status="unknown",iostat=ios)
25    write(*,*) "Testing matrix of size: "
26    do ii=1,grid_dim,1
27        call MatTest("time", grid(ii), "n", result)
28        write(*,*) floor(grid(ii))
29        results_all(ii,:) = result(:)
30        write(78,*) results_all(ii,:)
31    end do
32    write(*,*) "Done!"
33    close(78)
34
35 end program cos

```

To test the effective computing performance of my machine with different optimization flags, I wrote a Python script which has the following characteristics:

- it defines a grid of points which range with exponential spacing from N_{min} and N_{max} arbitrarily chosen; this is made to have equally spaced points on the log-scaled graphs that will be discussed below
- saves the grid on file
- it calls the compiler for `cos.f90`
- it repeats the operation for all optimization flags available
- it calls a `gnuplot` script to plot results.

The Python script `program.py` is the following:

```

1 import os
2 import numpy

```

```

3
4 fid=open("grid.dat","w+") #Open file
5 numbers=numpy.logspace(1.,numpy.log10(2500),num=10) # Create the grid
6 numbers=numpy.floor(numbers)
7 for item in numbers:
8     fid.write("%s\n" % item)
9 fid.close()
10
11 opt_flags = ["", "-O1", "-O2", "-O3", "-Ofast"]
12 for item in opt_flags: #cycle through opt flags
13     comp_comm="gfortran cos.f90 -o cos.out "+item
14     os.system(comp_comm) #Call the compiler
15     command = "./cos.out " + item
16     os.system(command) #Run the program
17
18 os.system("gnuplot plotres.gp") #Call gnuplot for graphs

```

This gnuplot script (plotres.gp) fits and plots all the computation times; this is done on different images for each optimization flag used. The results of the fits are displayed on terminal.

```

... line style definition ...
... set terminal and labels ...
f(x) = a*x**b # Fitting functions
g(x) = c*x**d
h(x) = e*x**p
...
fit f(x) "result.dat" using 1:2 via a,b
fit g(x) "result.dat" using 1:3 via c,d
# fit h(x) "result.dat" using 1:4 via e,p
plot "result.dat" using 1:2 title 'StdLoop' with points linestyle 1, \
    "result.dat" using 1:3 title 'InvLoop' with points linestyle 2, \
    "result.dat" using 1:4 title 'Intrinsic' with points linestyle 3, \
    f(x) with line linestyle 4, \
    g(x) with line linestyle 4, \
    # h(x) with line linestyle 4 # bad results! not plotted
... repeat for all opt flags ...

```

Results and self evaluation

Using Python script and `program cos` I tested the computation times for matrices with sizes spanning from 10 to 2500. I collected these data into graphs, which can be seen in figure 1 and figure 2. The first graph contains data collected from an unoptimized algorithm, while data on the remaining ones were collected from algorithms with increasing optimization.

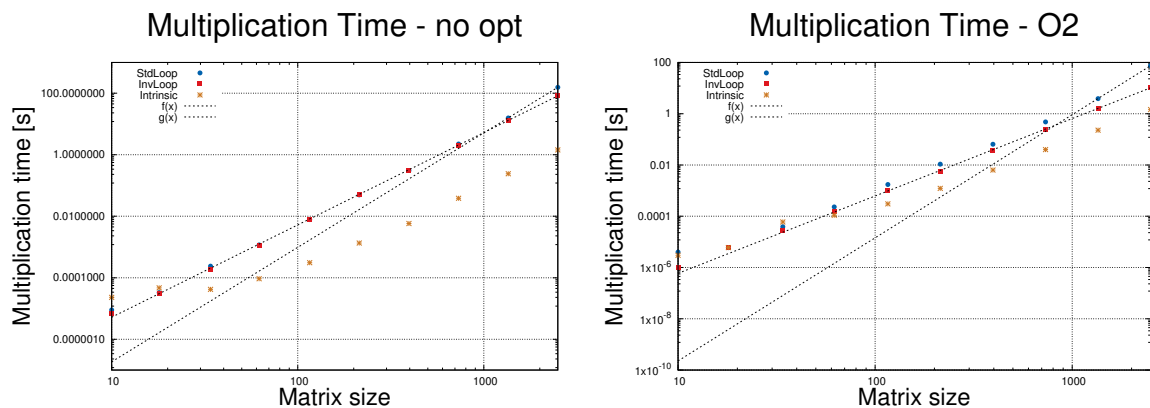


Figure 1: Computation time as a function of the matrix size, without opt flags and with -O2.

On graph I plotted power laws $f(x) = ax^b$, with a, b parameters retrieved from the data via fit; these power laws look like straight lines on these plots because of the double logarithmic scale. The

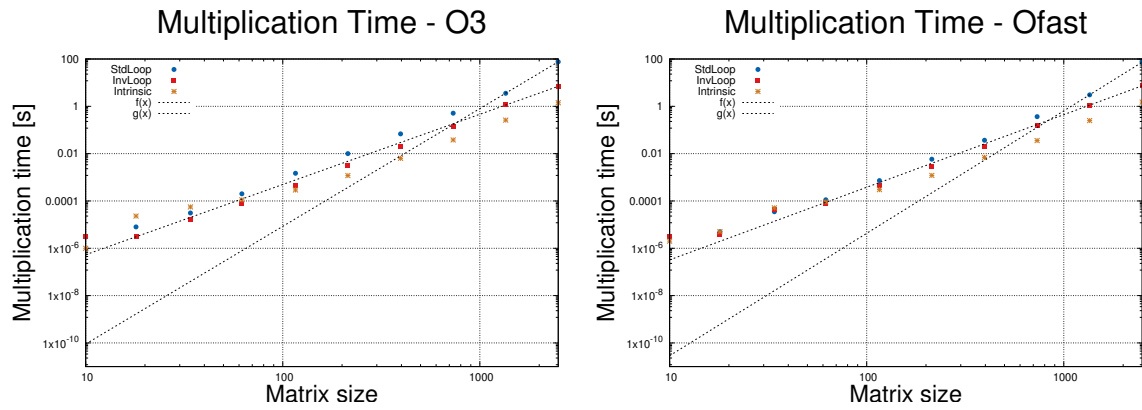


Figure 2: Computation time as a function of the matrix size, with -O3 and -Ofast.

agreement between data and the power law is good for the "InvLoop" method (non-optimized), while is poor for the others; fit results for "Intrinsic" multiplication method are not shown since this data present an evident nonlinear behavior in both log-log plots.

These graphs suggest that the computation times scale as power laws with the matrix size, at least for user-implemented methods. This statement might be verified building a more complex test involving repeated runs of the Python script, the elaboration of the data on a statistical base and the construction of confidence intervals for each data point.

Building this program I learned for the first time how to write a Python script and some ways to pass arguments between different languages.