

Multi-run script & Automated Fits

Umberto Maria Tomasini
Quantum Information Course

06/11/2018

Abstract

In this report we will take the code of the exercise 3 of the first week -which was about matrix- matrix multiplication via three different methods- and we will modify it in order to read the dimension of the matrix from a file. Then we will write a script which executes that program in a certain range of matrix's dimension, plotting the resulting Cpu_time over those dimensions. Lastly we will write another script which will automatically fit the data, compute and plot absolute residuals, relative residuals and their logarithm, via Gnuplot.

Theory

Code Development

Program Ex.3 modified

The code of the third exercise was modified as follows:

- the multiplication squared matrix- squared matrix via three different methods (row by row, column by column and by the intrinsic function *matmul*) is done for only one dimension. That dimension is taken as input by the file "MatDimension.txt".
- the three outputs (the cpu_times spent on each method) are printed in three different file: "Results-1" for the row by row method, "Results-2" for the column by column method and "Results-3" for the intrinsic function method (the access to file is "append"). The formatting of the output files was modified: `write(unit=unit,"(I4,4X,E16.9)", iostat=my_stat, iomsg=my_msg)`. The file is such that in every line there is an integer of max 4 figures (the matrix dimension) and another number which is written in scientific notation with nine significant figures (the cpu_time), separated by four spaces.
- In order to make more clear the execution of the script, we have chosen to turn the value of the variable `DEBUG` to `".false."` (the program was already executed without flaws).
- The computation of the total time spent on debugging is removed, as well as the printing of the instant of time.

The complete code is in the last section.

Scripting: making fits and plots

The following script executes the above program from a dimension `nmin` to a dimension `nmax`, with a certain step. The results are three file called "Results-1", "Results-2" and "Results-3". For every file of resulting data considered, the script gives as output:

- A linear fit, in a file called "linearfit-%d.pdf", where "%d" is the number of the "Results-" file considered. The fit is done taking the log of the resulting data. This is linear. In fact for high matrix dimension the computational cost of each method should be: $\Delta t \approx a \cdot n^3$, discarding the non-leading terms. Taking the logarithm: $\log(\Delta t) \approx 3 \log(n) + \log(a)$, which is linear in $\log(n)$.
- The absolute residuals $|\Delta f = f_{th} - f_{exp}|$, in a file called "absres-%d.pdf".
- The relative residuals $|\frac{\Delta f = f_{th} - f_{exp}}{f_{th}}|$, in a file called "relres-%d.pdf".
- The relative residuals $\log \left(\left| \frac{\Delta f = f_{th} - f_{exp}}{f_{th}} \right| \right)$, in a file called "logrelres-%d.pdf".
- The data of each fit, contained inside "parametersfit-%d."

Moreover all three sets of data are plotted on the same pdf, called "rawplot.pdf". The last output is another pdf, called "all3linearfit.pdf" which contains the plots of all three fitting functions (of the three methods). In the same figure the resulting data are plotted or, more precisely, their logarithm (in order to make the comparison with the fitting function reasonable). The parameters of those fits are contained in the file called "all3parametersfit.txt".

Listing 1: Scripting: making fits and plots

```
import numpy as np
import sys
import os
import subprocess

nmin=100
#lower bound dimension matrix
nmax=2525
#upper bound dimension matrix
stepp=25
#step of the increasing dimension

#####
#cycle to re-iterate the launch of the program
#that multiply two matrices in three different method
#and compute the computational cpu time needed
#with the dimension of the matrix that goes from
#nmin to nmax with step
#the result are saved in three files:
#"Results-1": multiplication row by row.
#"Results-2": multiplication by col by col.
#"Results-3": multiplication by intrinsic matmul.

#####

for ii in range(nmin,nmax,stepp):
f = open("MatDimension.txt","w")
#opening the file with the dimension of the matrix
f.write(str(ii))
f.close()
#write the dimension on the opened file
subprocess.call("./a.out")
#execute the fortran program

#####
#PLOTTING the three sets of data on the same pdf: "rawplot.
```

```

pdf"

#Within gnuplot the list of commands contained in "rawplot"
  is executed
#On the x-axis there is the dimension N of the matrix
#On the y-axis there is the Cpu time
#the x range is from nmin to nmax (to be inserted manually
  inside "rawplot")
#Results 1: by row. Results 2: by col. Results 3: intrinsic
  matmul.
#####

os.system("gnuplot_rawplot")
#Results 1: by row. Results 2: by col. Results 3: intrinsic
  matmul.

#####
#AUTOMATIC FITTING of the sets of data, separately.

#For every "Results-%d" file, the list of commands within "
  fitting" executes:
#linear fit (taking the log of the cubic law), inside "
  linearfit-%d.pdf"
#absolute residuals ( $\Delta f$ ), inside "absres-%d.pdf"
#relative residuals ( $\Delta f/f$ ), inside "relres-%d.pdf"
#logarithm of relative residuals ( $\log(\Delta f/f)$ ), inside "
  logrelres-%d.pdf"
#the x range is from nmin to nmax (to be inserted manually
  inside "rawplot")
#On the x-axis there is  $\log(N)$ , with N the dimension of the
  matrix
#The data of the fit are contained inside "parametersfit-%d."

#Logic of the script: the data contained in the considered "
  Results-%d" file
#are copied inside a temporary file "resultstemp". On that
  file are done the
#plots and the fit listed above. The results are stored in
  file with a
#temporary name, which are changed in the names listed above,
  which contain
#the number of the "Results-%d" file. This applies also for
  the data on the fit.

```

```
#This procedure is valid for a general number of "Results-%d"
file,
#contained in the same directory of the script.
#This number is saved in the file called "numresults.txt",
#from which it is read.
#####

os.system("find . -name \"Results-*\" | wc -l > numresults.
txt")
#prints the number of result files on a txt

num = open("numresults.txt", "r")
nn= int (num.readlines()[0])
num.close()
#read number of result files

for ii in range(1,nn+1):
#cycle over the "Results-%d" file
res = open("Results-%d" % ii, "r")
reslines = res.readlines()
#reading data from result file number ii

temp = open("resultstemp", "w")
for jj in reslines:
temp.write(jj)
#copying data in temporary file
res.close()
temp.close()
#closing files

os.system("gnuplot_fitting")
#linear fit of data, absolute residuals,
#relative residuals, log of realtive residuals

newnamelinear= ("linearfit-%d.pdf" %ii)
os.rename("templinearfit.pdf",newnamelinear)
#renaming the file of the linear fit

newnameabs= ("absres-%d.pdf" %ii)
os.rename("tempabsres.pdf",newnameabs)
#renaming the file of the absolute residuals

newnamerel= ("relres-%d.pdf" %ii)
os.rename("temprelres.pdf",newnamerel)
#renaming the file of the relative residuals
```

```
newnamelog= ("logrelres-%d.pdf" %ii)
os.rename("templogrelres.pdf",newnamelog)
#renaming the file of the log of the relative residuals

os.remove("resultstemp")
#removing the temporary file of results

fitlog = open("fit.log","r")
fitlines = fitlog.readlines()
#reading parameters of the fit from "fit.log"

fitlogd = open("parametersfit-%d" % ii,"w")
for jj in fitlines:
    fitlogd.write(jj)
#copying parameters on "parametersfit-%d"
fitlog.close()
fitlogd.close()
#closing files

os.remove("fit.log")
#removing "fit.log"

#####
#FITTING the three sets of data all together.

#linear fit (taking the log of the cubic law)
#Within gnuplot the list of commands contained in "
    all3fitting" is executed
#On the x-axis there is log(N), with N the dimension of the
    matrix
#On the y-axis there is log(Cpu time)
#the x range is from nmin to nmax (to be inserted manually
    inside "all3fitting")
#Results 1: by row. Results 2: by col. Results 3: intrinsic
    matmul.
#####

os.system("gnuplot_all3fitting")
#Results 1: by row. Results 2: by col. Results 3: intrinsic
    matmul.

flog = open("fit.log","r")
all3fitlines = flog.readlines()
```

```
#reading parameters of the fit from "fit.log"

all3fitlog = open("all3parametersfit", "w")
for jj in all3fitlines:
all3fitlog.write(jj)
#copying parameters on "all3parametersfit"

flog.close()
all3fitlog.close()
#closing files

os.remove("fit.log")
#removing "fit.log"

exit()
```

Gnuplot commands

In the script above are called via the command *os.system* three different lists of gnuplot commands.

Listing 2: Plotting the three sets of data together

```
set term pdf

set output "rawplot.pdf"

#name of the output file

set xr [100:3900]

#settin xrange from nmin to nmax

set xlabel "N"

set ylabel "Cpu_time_[sec]"

#on the x axis there is the dimension of the matrix
#on the y axis there is the cpu time

set title "Results-1:_by_row._Results-2:_by_col._Results-3:_\
intrinsic_matmul._"

set key top left

#the table of the contents is moved to the figure's top-left
```



```
plot "Results-1" u 1:2 w p lt rgb "blue", "Results-2" u 1:2 w
  p lt rgb "red", "Results-3" u 1:2 w p lt rgb "green"

#Results-1 (by row): blue
#Results-2 (by col): red
#Results-3 (by intrinsic): green

set term wxt

#set to normal terminal
```

Listing 3: Automatic fitting and residuals

```
#LINEAR FIT

set term pdf

set output "templinearfit.pdf"

#name of the output pdf

f(x)= a+b*x

fit f(x) "resultstemp" u (log($1)):(log($2)) via a,b

#fitting the data taking the logarithm of both columns
#via a linear function

set xr [log(100):log(3900)]

#setting xrange

set xlabel "log(N) "

set ylabel "log(Cpu_time) "

plot "resultstemp" u (log($1)):(log($2)) w p lt rgb "blue"
  notitle, f(x) notitle

#plotting of the data (the log(data)) and
#the interpoling function

#ABSOLUTE RESIDUALS
```

```
set output "tempabsres.pdf"

#name of the output pdf

p(x,y)=abs(f(x)-y)

#this function computes absolute residuals

set xr [log(100):log(3900)]

set xlabel "log(N)"

set ylabel "|{/Symbol_D}f|" enhanced

plot "resultstemp" u (log($1)): (p(log($1),log($2))) w p lt
    rgb "green" notitle
#the p function accept as input the two columns of data

#RELATIVE RESIDUALS

set output "temprelres.pdf"

#name of the output pdf

p(x,y)=abs((f(x)-y)/f(x))

#this function computes relative residuals

set xr [log(100):log(3900)]

set xlabel "log(N)"

set ylabel "|{/Symbol_D}f/f|" enhanced

plot "resultstemp" u (log($1)): (p(log($1),log($2))) w p lt
    rgb "brown" notitle
#the p function accept as input the two columns of data

#LOGARITHM OF RELATIVE RESIDUALS

set output "templogrelres.pdf"

#name of the output pdf

p(x,y)=log(abs((f(x)-y)/f(x)))
```

```

#this function is the logarithm of relative residuals

set xr [log(100):log(3900)]

set xlabel "log(N) "

set ylabel "log(|{/Symbol_D}f/f|) " enhanced

plot "resultstemp" u (log($1)): (p(log($1),log($2))) w p lt
    rgb "red" notitle
#the p function accept as input the two columns of data

set term wxt

```

Listing 4: Plotting the three fits together

```

set term pdf

set output "all3linearfit.pdf"

#setting the output pdf

f(x)= a+b*x

g(x)= c+d*x

h(x)= e+f*x

#three different linear function

fit f(x) "Results-1" u (log($1)): (log($2)) via a,b

fit g(x) "Results-2" u (log($1)): (log($2)) via c,d

fit h(x) "Results-3" u (log($1)): (log($2)) via e,f

#three different linear fits

set xr [log(100):log(3900)]

set xlabel "log(N) "

set ylabel "log(Cpu_time) "

```

```

set title "Results-1:_by_row._Results-2:_by_col._Results-3:_
    intrinsic_matmul._"

set key top left

#the table of the contents is moved to the figure's top-left

plot "Results-1" u (log($1)):(log($2)) w p lt rgb "blue", f(x)
    lt rgb "blue" notitle, "Results-2" u (log($1)):(log($2))
    w p lt rgb "red", g(x) lt rgb "red" notitle, "Results-3"
    u (log($1)):(log($2)) w p lt rgb "green", h(x) lt rgb "
    green" notitle

#Results-1 and fitting f(x) (by row): blue
#Results-2 and fitting g(x) (by col): red
#Results-3 and fitting h(x) (by intrinsic): green

set term wxt

```

Results

Estimating N_{max}

In order to decide the range of the matrices dimension (from n_{min} to n_{max}) e the step of the increasing dimension, it was decided firstly to go from 100 to 2525 with $step=25$. The execution took more or less three hours and an half. The parameters of the three fits obtained were:

	Intercept	Slope
By row	-23.6514	3.75172
By col	-23.6811	3.7648
By matmul	-21.1094	2.95827

Using those parameters, the following Python code was implemented. Its goal is to estimate the time needed for further computations.

Listing 5: Estimating the total cpu time

```

import numpy as np

nmin=2600
nmax=4000
step=100
accum=0

```

```

for ii in range(nmin,nmax,step):
accum= accum+ np.exp(-23.6514)*(ii)**3.75172 + np.exp
      (-23.6811)*(ii)**3.7648 +np.exp(-21.1094)*(ii)**2.95827
#the parameters of the fits are used

print(accum)
#number of seconds
print(accum/3600)
#number of hours

```

Obviously this program takes into account only the contribution $\approx O(n^3)$ of the multiplications, disregarding the minor contributions of the rest of the program. Deciding a priori that the maximum range of time given to further computations would be 7 hours/ 7 hours and half (a night's sleep), the range was set from 2600 to 4000 with a step of 100 (the 4000 was not computed).

Combining the results obtained from 100 to 2525 with a step of 25 and from 2600 to 4000 with a step of 100, the results are the following.

Results from 100 to 3900

Plot of the three sets of data




Figure 1: Plot of the three sets of data


Fit and residuals: row by row method

Figure 2: Linear fit: row by row




Figure 3: Absolute residuals: row by row




Figure 4: Relative residuals: row by row




Figure 5: Log of relative residuals: row by row


Fit and residuals: column by column method

Figure 6: Linear fit: column by column




Figure 7: Absolute residuals: column by column




Figure 8: Relative residuals: column by column




Figure 9: Log of relative residuals: column by column


Fit and residuals: intrinsic function method

Figure 10: Linear fit: intrinsic function




Figure 11: Absolute residuals: intrinsic function




Figure 12: Relative residuals: intrinsic function




Figure 13: Log of relative residuals: intrinsic function

Plot of the three fitting function

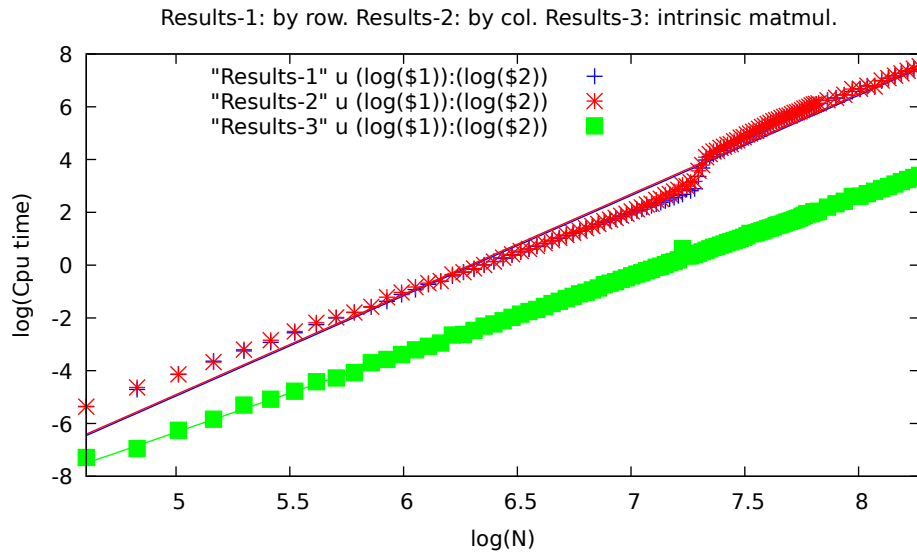


Figure 14: Plot of the three fitting function

Parameters of the three fits (the complete fit.log is below).

	Intercept	Slope
By row	-23.9417	3.797
By col	-23.8943	3.79811
By matmul	-21.1703	2.96775

Motivated by the fact that in the linear fits of the first two methods two different linear trends can be recognized, we repeated the analysis on a first zone from 100 to 1400 and on a second zone from 1800 to 3900. For sake of completeness we also repeated the analysis on the third method, although there is only one main trend.

Results from 100 to 1400

Plot of the three sets of data. From 100 to 1400

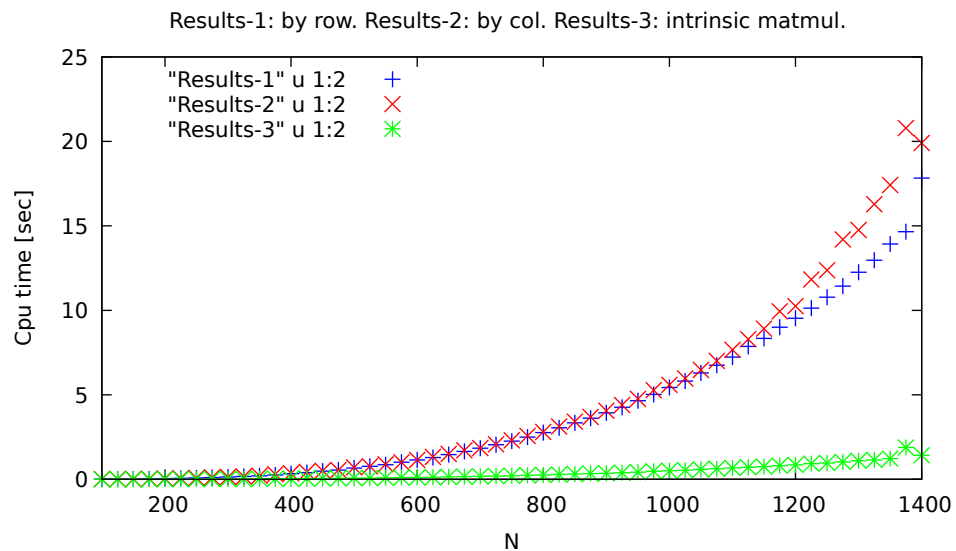


Figure 15: Plot of the three sets of data. From 100 to 1400

Fit and residuals: row by row method. From 100 to 1400




Figure 16: Linear fit: row by row. From 100 to 1400




Figure 17: Absolute residuals: row by row. From 100 to 1400




Figure 18: Relative residuals: row by row. From 100 to 1400




Figure 19: Log of relative residuals: row by row. From 100 to 1400

Fit and residuals: column by column method. From 100 to 1400




Figure 20: Linear fit: column by column. From 100 to 1400




Figure 21: Absolute residuals: column by column. From 100 to 1400




Figure 22: Relative residuals: column by column. From 100 to 1400n




Figure 23: Log of relative residuals: column by column. From 100 to 1400

Fit and residuals: intrinsic function method. From 100 to 1400




Figure 24: Linear fit: intrinsic function. From 100 to 1400




Figure 25: Absolute residuals: intrinsic function. From 100 to 1400




Figure 26: Relative residuals: intrinsic function. From 100 to 1400




Figure 27: Log of relative residuals: intrinsic function. From 100 to 1400

Plot of the three fitting function. From 100 to 1400

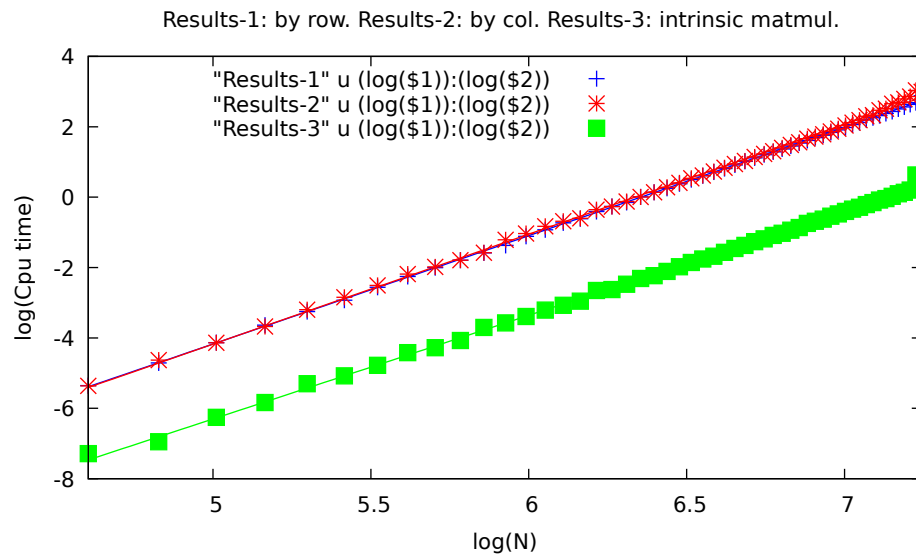


Figure 28: Plot of the three fitting function. From 100 to 1400

Parameters of the three fits (from 100 to 1400 (matrix dim)). The complete data are in the below subsection.

	Intercept	Slope
By row	-19.5703	3.07942
By col	-19.7937	3.1236
By matmul	-21.0333	2.94588

Results from 1800 to 3900

Plot of the three sets of data. From 1800 to 3900

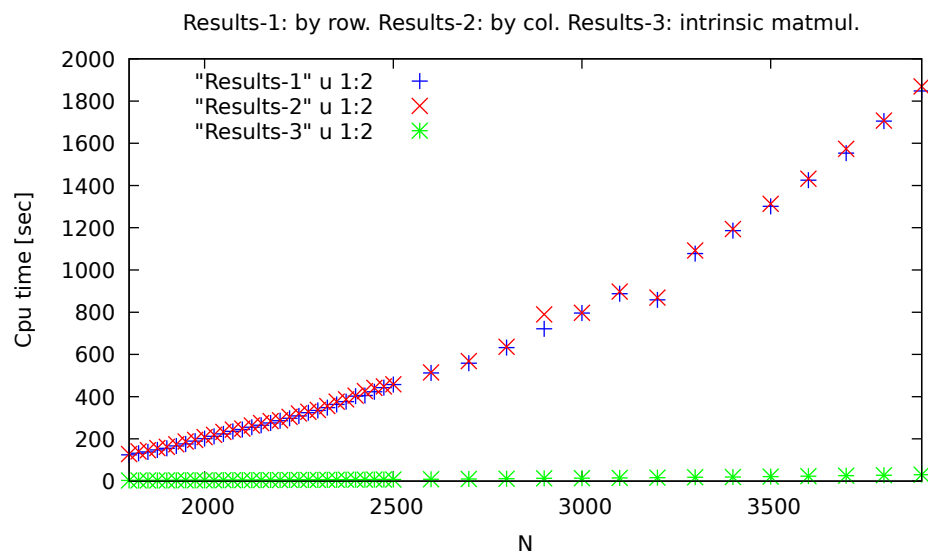


Figure 29: Plot of the three sets of data. From 1800 to 3900


Fit and residuals: row by row method. From 1800 to 3900

Figure 30: Linear fit: row by row. From 1800 to 3900




Figure 31: Absolute residuals: row by row. From 1800 to 3900




Figure 32: Relative residuals: row by row. From 1800 to 3900




Figure 33: Log of relative residuals: row by row. From 1800 to 3900


Fit and residuals: column by column method. From 1800 to 3900

Figure 34: Linear fit: column by column. From 1800 to 3900




Figure 35: Absolute residuals: column by column. From 1800 to 3900




Figure 36: Relative residuals: column by column. From 1800 to 3900




Figure 37: Log of relative residuals: column by column. From 1800 to 3900


Fit and residuals: intrinsic function method. From 1800 to 3900

Figure 38: Linear fit: intrinsic function. From 1800 to 3900




Figure 39: Absolute residuals: intrinsic function. From 1800 to 3900




Figure 40: Relative residuals: intrinsic function. From 1800 to 3900




Figure 41: Log of relative residuals: intrinsic function. From 1800 to 3900

Plot of the three fitting function. From 1800 to 3900




Figure 42: Plot of the three fitting function. From 1800 to 3900

Parameters of the three fits (from 1800 to 3900 (matrix dim)). The complete data are in the below subsection.

	Intercept	Slope
By row	-20.4728	3.39082
By col	-20.0473	3.33968
By matmul	-21.6692	3.03309

Self-Evaluation

Comments on the Results

The most striking features of the plot of the three sets of data is the similarity between the first and the second method (this is reasonable: the two methods are very similar) and the difference between those two and the third method, which is much faster than the other.

The liner fits of the first two methods on the whole range [100,3900] present the following trend: until some point between 7 and 7.5 (1100 and 1800 in matrix dimension) there is one linear trend, then there is a transient and then another linear trend, from 7.5 (more or less). This observation is motivated by the fact that the absolute residuals, the relative residuals and the logarithm of relative residuals present a discontinuity in the so called transient zone. Moreover, the absolute residuals in the zone from 100 to the transient zone suggests that the points are aligned along a line which is incident with the fitting line, with the intersection around 6 (400 in matrix dimension, more or less). After the transient there is another

descent, which recalls the one in the first zone. As a consequence, we can infer there is another alignment of the point in that high-dimension zone. As a consequence of those alignments, the slopes are very different from the theoretical ones: around 3.8.

On the other hand, the analysis on the third method presents a good linear fit, supported by low absolute residuals, relative residuals and logarithms of relative residuals.

In all the three methods, in the relative residuals and in the logarithm of relative residuals we can see a small zone which is much higher than the others. In that small zone the fitting function is near to 0, so the relative residuals and their logarithm are very high, respect to the others.

The analysis between 100 and 1400 shows us a good linear fit, for all the three methods, supported by low absolute residuals, relative residuals and logarithms of relative residuals. It is noticeable that the slopes are closer to the theoretical 3 respect to the ones of the overall fit.

In the absolute residuals, the relative residuals and the logarithm of relative residuals of the fits from 1800 to 3900 we can see a discontinuity between 7.8 and 7.9. This is due to the fact that before 2500 the steps between dimensions was 25, then 100. As a consequence, there are a lot more points before 2500. Nevertheless, the residuals are low and therefore the fit is good. On the other hand, the obtained slopes are different from the expected ones for the first two methods: 3.3 or 3.4 instead of 3.

What we have learnt:

- Writing a Python Script which communicates automatically with the terminal and Gnuplot, in addition to automatically executing programs.
- Performing via Gnuplot linear fits, plotting absolute residuals, relative residuals and their logarithm.
- Commenting and thinking on the results.
- Estimating the time needed for high-cost computations.

What can be done in the future:

- Understand why the points align along two different lines, one in the low-dimension zone and one in the high-dimension zone.

Data of the three fits: from 100 to 3900

Mon Oct 29 09:45:06 2018

FIT: data read from "Results-1" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 111

residuals are weighted equally (unit weight)

function used for fitting: f(x)

f(x)= a+b*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	a	b
0	3.4755286224e+03	0.00e+00	5.10e+00	1.000000e+00	1.000000e+00
4	2.2068480154e+01	-3.12e-01	5.10e-04	-2.394167e+01	3.796996e+00

After 4 iterations the fit converged.

final sum of squares of residuals : 22.0685

rel. change during last iteration : -3.12165e-06

degrees of freedom	(FIT_NDF)	:	109
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.449959
variance of residuals	(reduced chisquare) = WSSR/ndf	:	0.202463

Final set of parameters	Asymptotic Standard Error
=====	=====
a = -23.9417	+/- 0.3789 (1.583%)
b = 3.797	+/- 0.05308 (1.398%)

correlation matrix of the fit parameters:

	a	b
a	1.000	
b	-0.994	1.000

Mon Oct 29 09:45:06 2018

FIT: data read from "Results-2" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 111

residuals are weighted equally (unit weight)

function used for fitting: g(x)

$g(x) = c + d \cdot x$

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	c	d
0	3.4109617706e+03	0.00e+00	5.10e+00	1.000000e+00	1.000000e+00
4	1.9350825095e+01	-3.55e-01	5.10e-04	-2.389430e+01	3.798111e+00

After 4 iterations the fit converged.

final sum of squares of residuals : 19.3508

rel. change during last iteration : -3.5468e-06

degrees of freedom	(FIT_NDF)	:	109
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.421344
variance of residuals	(reduced chisquare) = WSSR/ndf	:	0.177531

Final set of parameters		Asymptotic Standard Error	
=====		=====	
c	= -23.8943	+/- 0.3548	(1.485%)
d	= 3.79811	+/- 0.0497	(1.309%)

correlation matrix of the fit parameters:

	c	d
c	1.000	
d	-0.994	1.000

Mon Oct 29 09:45:06 2018

FIT: data read from "Results-3" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 111

residuals are weighted equally (unit weight)

function used for fitting: $h(x)$

$h(x) = e + f \cdot x$

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	e	f
0	7.7673995350e+03	0.00e+00	5.10e+00	1.000000e+00	1.000000e+00
5	3.1049987943e-01	-5.61e-09	5.10e-05	-2.117031e+01	2.967751e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.3105

rel. change during last iteration : -5.61369e-14

degrees of freedom (FIT_NDF) : 109
 rms of residuals (FIT_STDFIT) = sqrt(WSSR/ndf) : 0.0533725
 variance of residuals (reduced chisquare) = WSSR/ndf : 0.00284862

Final set of parameters	Asymptotic Standard Error
=====	=====
e = -21.1703	+/- 0.04494 (0.2123%)
f = 2.96775	+/- 0.006296 (0.2122%)

correlation matrix of the fit parameters:

e	f	
e	1.000	
f	-0.994	1.000

Data of the three fits: from 100 to 1400

 Wed Oct 31 09:33:48 2018

FIT: data read from "Results-1" u (log(\$1)):(log(\$2))
 format = x:z
 #datapoints = 53
 residuals are weighted equally (unit weight)

function used for fitting: f(x)
 f(x)= a+b*x
 fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	a	b
0	2.8403361635e+03	0.00e+00	4.63e+00	1.000000e+00	1.000000e+00
5	3.5812761533e-02	-1.02e-06	4.63e-05	-1.957031e+01	3.079415e+00

After 5 iterations the fit converged.
 final sum of squares of residuals : 0.0358128
 rel. change during last iteration : -1.02136e-11

degrees of freedom (FIT_NDF) : 51
 rms of residuals (FIT_STDFIT) = sqrt(WSSR/ndf) : 0.0264993
 variance of residuals (reduced chisquare) = WSSR/ndf : 0.000702211

Final set of parameters		Asymptotic Standard Error	
=====		=====	
a	= -19.5703	+/- 0.03495	(0.1786%)
b	= 3.07942	+/- 0.005399	(0.1753%)

correlation matrix of the fit parameters:

a	b	
a		1.000
b	-0.995	1.000

Wed Oct 31 09:33:48 2018

FIT: data read from "Results-2" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 53

residuals are weighted equally (unit weight)

function used for fitting: g(x)

g(x)= c+d*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	c	d
0	2.7987669617e+03	0.00e+00	4.63e+00	1.0000000e+00	1.0000000e+00
5	2.5970853675e-01	-1.44e-07	4.63e-05	-1.979373e+01	3.123602e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.259709

rel. change during last iteration : -1.44256e-12

degrees of freedom	(FIT_NDF)	:	51
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.0713605
variance of residuals (reduced chisquare)	= WSSR/ndf	:	0.00509232

Final set of parameters		Asymptotic Standard Error	
=====		=====	
c	= -19.7937	+/- 0.09411	(0.4754%)
d	= 3.1236	+/- 0.01454	(0.4655%)

correlation matrix of the fit parameters:

c	d	
c		1.000

d -0.995 1.000

Wed Oct 31 09:33:48 2018

FIT: data read from "Results-3" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 53

residuals are weighted equally (unit weight)

function used for fitting: h(x)

h(x)= e+f*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	e	f
0	4.8823868816e+03	0.00e+00	4.63e+00	1.000000e+00	1.000000e+00
5	2.5916307452e-01	-1.61e-07	4.63e-05	-2.103325e+01	2.945878e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.259163

rel. change during last iteration : -1.60795e-12

degrees of freedom	(FIT_NDF)	:	51
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.0712855
variance of residuals	(reduced chisquare) = WSSR/ndf	:	0.00508163

Final set of parameters	Asymptotic Standard Error
=====	=====
e = -21.0333	+/- 0.09401 (0.447%)
f = 2.94588	+/- 0.01453 (0.4931%)

correlation matrix of the fit parameters:

e	f
e	1.000
f	-0.995 1.000

Data of the three fits: from 1800 to 3900

Wed Oct 31 17:56:51 2018

FIT: data read from "Results-1" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 43

residuals are weighted equally (unit weight)

function used for fitting: f(x)

f(x)= a+b*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	a	b
0	3.5393020278e+02	0.00e+00	5.56e+00	1.000000e+00	1.000000e+00
5	1.2662992118e-01	-6.01e-02	5.56e-05	-2.047284e+01	3.390820e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.12663

rel. change during last iteration : -6.01076e-07

degrees of freedom	(FIT_NDF)	:	41
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.0555746
variance of residuals	(reduced chisquare) = WSSR/ndf	:	0.00308853

Final set of parameters	Asymptotic Standard Error
=====	=====
a = -20.4728	+/- 0.3003 (1.467%)
b = 3.39082	+/- 0.03847 (1.135%)

correlation matrix of the fit parameters:

a	b
a	1.000
b	-1.000 1.000

Wed Oct 31 17:56:51 2018

FIT: data read from "Results-2" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 43

residuals are weighted equally (unit weight)

function used for fitting: g(x)

g(x)= c+d*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	c	d
0	3.4701109840e+02	0.00e+00	5.56e+00	1.0000000e+00	1.0000000e+00
5	1.1751538879e-01	-6.22e-02	5.56e-05	-2.004734e+01	3.339684e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.117515

rel. change during last iteration : -6.22253e-07

degrees of freedom	(FIT_NDF)	:	41
rms of residuals	(FIT_STDFIT) = sqrt(WSSR/ndf)	:	0.0535372
variance of residuals (reduced chisquare)	= WSSR/ndf	:	0.00286623

Final set of parameters	Asymptotic Standard Error
=====	=====
c = -20.0473	+/- 0.2893 (1.443%)
d = 3.33968	+/- 0.03706 (1.11%)

correlation matrix of the fit parameters:

c	d
c	1.000
d	-1.000 1.000

Wed Oct 31 17:56:51 2018

FIT: data read from "Results-3" u (log(\$1)):(log(\$2))

format = x:z

#datapoints = 43

residuals are weighted equally (unit weight)

function used for fitting: h(x)

h(x)= e+f*x

fitted parameters initialized with current variable values

iter	chisq	delta/lim	lambda	e	f
0	2.0010699954e+03	0.00e+00	5.56e+00	1.0000000e+00	1.0000000e+00
5	1.9261685541e-02	-4.38e-01	5.56e-05	-2.166919e+01	3.033094e+00

After 5 iterations the fit converged.

final sum of squares of residuals : 0.0192617

rel. change during last iteration : -4.38012e-06

```

degrees of freedom      (FIT_NDF)                      : 41
rms of residuals        (FIT_STDFIT) = sqrt(WSSR/ndf)   : 0.0216748
variance of residuals   (reduced chisquare) = WSSR/ndf : 0.000469797

```

Final set of parameters	Asymptotic Standard Error
=====	=====
e = -21.6692	+/- 0.1171 (0.5405%)
f = 3.03309	+/- 0.015 (0.4947%)

correlation matrix of the fit parameters:

e	f
e	1.000
f	-1.000 1.000

Complete Code of the Program

Listing 6: Fortran code Ex.3 modified

```

\brieff \b Multiplication matrix-matrix and its
      computational time
C>
C> ===== DOCUMENTATION =====
C> Definition:
C> =====
C>
C>      SUBROUTINE ROWROWMATMUL (nn,m1,m2,mris1,
      timetot1)
C>
C>      .. Scalar Arguments ..
C>      INTEGER*2          nn
C>      REAL              timetot1
C>      ..
C>      .. Array Arguments ..
C>      INTEGER*2          m1(nn,nn)
C>      INTEGER*2          m2(nn,nn)
C>      INTEGER*2          mris1(nn,nn)
C>      ..
C>
C> Purpose
C> =====
C>
C> \details \b Purpose:
C> \verbatim

```

```

C>
C> ROWROWMATMUL does a matrix-matrix multiplication
    between two matrices
C> given as input: m1 and m2. The resulting matrix is
    mris1. ROWROWMATMUL
C> deals with INTEGER*2 squared matrices, with
    numbers of rows and columns equal C> to nn. In the
    entry (i,j) of mris1 there is the scalar product
    between the C> i-th row of m1 and the j-th column
    of m2.
C> The multiplication is row by row because while i
    is fixed j
C> goes from 1 to nn. In other words the j cycle is
    inside the i cycle.
C> Calling the function CPU_TIME at the beginning of
    that multiplication, we assign the initial time to
    REAL variable start1. The same is done for the
    ending C> time, in REAL variable finish1. The
    total time timetot1 is the difference between the
    finish time and the ending time.

C>\endverbatim
C>
C> Arguments:
C> =====
C>

C> \param[in] nn
C> \verbatim
C>         nn is INTEGER*2
C>         The dimension of the squared array m1, m2
            and mris1
C> \endverbatim
C> \param[in] m1
C> \verbatim
C>         m1 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[in] m2
C> \verbatim
C>         m2 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[out] m2
C> \verbatim

```

```

C>          mris1 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C>
C> \param[out] timetot1
C> \verbatim
C>          timetot1 is REAL
C> \endverbatim
C>
C> Authors:
C> =====
C>
C> \author Univ. of Padua
C>
C> \date 23 October 2018
C>
C>
C>          SUBROUTINE COLCOLMATMUL (nn,m1,m2,mris2,
    timetot2)
C>
C>          .. Scalar Arguments ..
C>          INTEGER*2          nn
C>          REAL               timetot2
C>          ..
C>          .. Array Arguments ..
C>          INTEGER*2          m1 (nn,nn)
C>          INTEGER*2          m2 (nn,nn)
C>          INTEGER*2          mris2 (nn,nn)
C>          ..
C>
C> Purpose
C> =====
C>
C>\details \b Purpose:
C>\verbatim
C>
C> COLCOLMATMUL does a matrix-matrix multiplication
    between two matrices
C> given as input: m1 and m2. The resulting matrix is
    mris2. COLCOLMATMUL
C> deals with INTEGER*2 squared matrices, with
    numbers of rows and columns equal C> to nn. In the
    entry (j,i) of mris1 there is the scalar product
    between the j-th row of m1 and the i-th column of
    m2.

```

```
C> The multiplication is row by row because while i
    is fixed j
C> goes from 1 to nn. In other words the j cycle is
    inside the i cycle.
C> Calling the function CPU_TIME at the beginning of
    that multiplication, we assign the initial time to
    REAL variable start2. The same is done for the
    ending time, in REAL variable finish2. The total
    time timetot2 is the difference between the finish
    time and the ending time.
```

```
C>\endverbatim
```

```
C>
```

```
C> Arguments:
```

```
C> =====
```

```
C>
```

```
C> \param[in] nn
```

```
C> \verbatim
```

```
C>         nn is INTEGER*2
```

```
C>         The dimension of the squared array m1, m2
    and mris2
```

```
C> \endverbatim
```

```
C> \param[in] m1
```

```
C> \verbatim
```

```
C>         m1 is INTEGER*2 ARRAY (nn,nn)
```

```
C> \endverbatim
```

```
C>
```

```
C> \param[in] m2
```

```
C> \verbatim
```

```
C>         m2 is INTEGER*2 ARRAY (nn,nn)
```

```
C> \endverbatim
```

```
C>
```

```
C> \param[out] m2
```

```
C> \verbatim
```

```
C>         mris2 is INTEGER*2 ARRAY (nn,nn)
```

```
C> \endverbatim
```

```
C>
```

```
C>
```

```
C> \param[out] timetot2
```

```
C> \verbatim
```

```
C>         timetot2 is REAL
```

```
C> \endverbatim
```

```
C>
```

```
C> Authors:
```

```

C> =====
C>
C> \author Univ. of Padua
C>
C> \date 23 October 2018
C>
C>
C>      SUBROUTINE MATMULINTRINSIC (nn,m1,m2,mris1,
    timetot1)
C>
C>      .. Scalar Arguments ..
C>      INTEGER*2          nn
C>      REAL              timetot3
C>      ..
C>      .. Array Arguments ..
C>      INTEGER*2          m1(nn,nn)
C>      INTEGER*2          m2(nn,nn)
C>      INTEGER*2          mris3(nn,nn)
C>      ..
C>
C> Purpose
C> =====
C>
C>\details \b Purpose:
C>\verbatim
C>
C> MATMULINTRINSIC does a matrix-matrix
    multiplication between two matrices
C> given as input: m1 and m2. The resulting matrix is
    mris1. MATMULINTRINSIC
C> deals with INTEGER*2 squared matrices, with
    numbers of rows and columns equal C> to nn.
C> The multiplication is done via the intrinsic
    function mris3=matmul(m1,m2).
C> Calling the function CPU_TIME at the beginning of
    that multiplication, we assign the initial time
    to REAL variable start3. The same is done for the
    ending C> time, in REAL variable finish3. The
    total time timetot3 is the difference between the
    finish time and the ending time.

C>\endverbatim
C>
C> Arguments:
C> =====

```



```
C>

C> \param[in] nn
C> \verbatim
C>         nn*2 is INTEGER
C>         The dimension of the squared array m1, m2
C>         and mris3
C> \endverbatim
C> \param[in] m1
C> \verbatim
C>         m1 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[in] m2
C> \verbatim
C>         m2 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[out] m2
C> \verbatim
C>         mris3 is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C>
C> \param[out] timetot1
C> \verbatim
C>         timetot3 is REAL
C> \endverbatim
C>
C> Authors:
C> =====
C>
C> \author Univ. of Padua
C>
C> \date 23 October 2018
C>
C>         SUBROUTINE CHECKDIM(nn,nncheck,debug)
C>
C>         .. Scalar Arguments ..
C>         INTEGER*2          nn
C>         INTEGER*2          nncheck
C>         LOGICAL             debug
C>         ..
C>         .. Array Arguments ..
C>         INTEGER*2          m(nn,nn)
```

```
C>          INTEGER*2          mcheck(nn,nn)
C>          ..
C> Purpose
C> =====
C>
C>\details \b Purpose:
C>\verbatim
C>
C> CHECKDIM checks if the dimension of the matrix are
    correct. Firstly checks
C> if the the dimension INTEGER*2 is above 10000. If
    it is so, it prints a WARNING
C> message because it takes too much time. Secondly,
    it checks if nn is inferior
C> to 1, printing a WARNING message. Lastly, it
    checks if nn is actually nncheck
C> as it should be. If it is not true, it prints a
    WARNING message, stating which
C> should be the dimension and the actual one. If
    everything goes well, it prints "okay:_right_
    dimensions_matrices",
C> and the actual dimension.

C>\endverbatim
C>
C> Arguments:
C> =====
C>
C> \param[in] nn
C> \verbatim
C>          nn is INTEGER*2
C>          The dimension of the squared arrays
C> \endverbatim
C> \param[in] qq
C> \verbatim
C>          qq is INTEGER*2
C>          the number of the cycle, it should be nn=
    qq*100
C> \endverbatim
C>
C> \param[inout] debug
C> \verbatim
C>          debug is LOGICAL
C> \endverbatim
```

```

C>
C> Authors:
C> =====
C>
C> \author Univ. of Padua
C>
C> \date 23 October 2018

C>      SUBROUTINE CHECKMAT (nn,m,mcheck, debug)
C>
C>      .. Scalar Arguments ..
C>      INTEGER*2          nn
C>      LOGICAL            debug
C>      ..
C>      .. Array Arguments ..
C>      INTEGER*2          m(nn,nn)
C>      INTEGER*2          mcheck (nn,nn)
C>      ..
C> Purpose
C>
C> Purpose
C> =====
C>
C>\details \b Purpose:
C>\verbatim
C>
C> CHECKMAT checks if the two input matrices m and
C> mcheck are equal. This is done
C> via a do cycle which checks the equality between
C> the two matrices element
C> by element. For each entry which is not equal, the
C> internal variable INTEGER*4 accum
C> increas by 1 (starting from 0). If accum in the
C> end is bigger than 0,
C> a WARNING message states how much entries between
C> the two matrices are
C> different. Otherwise, it is printed "Same_matrices
C> ".
C>\endverbatim
C>
C> Arguments:
C> =====
C>
C> \param[in] nn

```

```
C> \verbatim
C>           nn is INTEGER*2
C>           The dimension of the squared arrays
C> \endverbatim
C>
C> \param[in]
C> \verbatim
C>           m is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[in]
C> \verbatim
C>           mcheck is INTEGER*2 ARRAY (nn,nn)
C> \endverbatim
C>
C> \param[inout] debug
C> \verbatim
C>           debug is LOGICAL
C> \endverbatim
C>
C> Authors:
C> =====
C>
C> \author Univ. of Padua
C>
C> \date 23 October 2018

C>
C>           PROGRAM test_performance_mulmat

C> Purpose
C> =====
C>
C>\details \b Purpose:
C>\verbatim
C>
C>This program implements matrix-matrix
multiplication in three different method: row by
row, column by column and by intrinsic function.
For each method an apposite subroutine is called:
ROWROWMATMUL, COLCOLMATMUL and MATMULINTRINSIC.
For each method the CPU_TIME is computed. This is
done for squared matrix, which dimension is taken
as input from a file, called "MatDimension.txt".
The first matrix is such that at the entry (i,j)
```

there is the value $i+j$, whilst the second matrix is such that at the entry (i,j) there is the value $i*j$. At the beginning, the matrices are allocated and in the ending are deallocated. The results are printed on three file, one for each method, called "Results-1" for the row by row method, "Results-2" for the col by col method and "Results-3" for the intrinsic function method. The format is the following: 4 figures for one integer, four spaces, and lastly a real number written in scientific notation, with nine significant figures.

C>

C> DEBUGGING of the program. This is controlled with a CHARACTER*1 variable called CHOICE. The programmer can choose to give the choice to do or not the debugging to the user keeping or changing the value of CHOICE. Keeping its default value "X", the user can choose at the beginning of the program to do the debug, via a printed choice on the terminal.

C> If he says no, the logical variable DEBUG turns .FALSE. and all the debugging routines do not work.

C> If he says yes, the logical variable DEBUG turns .TRUE. and therefore all the debugging procedures below are executed. If one of the them reports an error, a warning message is printed.

C> Otherwise, if the programmer changes the value of CHOICE in "n", the user has not this choice and the debugging is not done (DEBUG is .FALSE.). Lastly, changing the value of CHOICE in "y", the user has not the choice and the debugging is done (DEBUG is .TRUE.).

C> The debugging is done printing various checkpoints throughout the program, in specific points: at the beginning of each cycle, initialization of the input matrices, calling the subroutines of the three methods, checking the resulting matrices.

C> Moreover the subroutines CHECKDIM and CHECKMAT are used a lot, interfaced via the operator CHECK. CHECKDIM is used every time a subroutine is called. CHECKMAT checks after the calling of the three method that the input matrices are still the same (after each one). Moreover, it checks if the resulting matrices mris1, mris and mris3 are the

same at the `end` of each `cycle`.

C> Lastly, error handling commands are added for the functions `OPEN`, `WRITE`, `ALLOCATE` and `DEALLOCATE`. They `print` which error the functions reports and its relative error message.

C> If one of the these procedures reports an error, a warning message is printed.

```
C>\endverbatim
C>  Authors:
C>  =====
C>
C>  \author Univ. of Padua
C>
C>  \date 23 October 2018
```

```
C>
```

```
=====

!MODULE MAT-MULTIPLICATIONS
module matmultiplications

contains
subroutine rowrowmatmul(nn,m1,m2,mris1,timetot1)
!first method:row by row
implicit none
integer*2 ii, jj, kk, nn
integer*2, dimension(nn,nn) :: m1
integer*2, dimension(nn,nn) :: m2
integer*2, dimension(nn,nn):: mris1
real :: finish1, start1,timetot1
call cpu_time (start1)!initial time
do ii=1,nn
do jj=1,nn
do kk=1,nn
mris1(ii,jj)=mris1(ii,jj)+m1(ii,kk)*m2(kk,jj)
!SCALAR PRODUCT BTW
!ii-th ROW OF m1 AND j-th COLUMN OF m2
!IT GOES IN THE ENTRY (i,j) OF mris1
end do
end do
end do
call cpu_time (finish1)!finish time
timetot1=finish1-start1
end subroutine rowrowmatmul
```

```
subroutine colcolmatmul(nn,m1,m2,mris2,timetot2)
!second method:col by col
implicit none
integer*2 ii, jj, kk, nn
integer*2, dimension(nn,nn) :: m1
integer*2, dimension(nn,nn) :: m2
integer*2, dimension(nn,nn) :: mris2
real :: finish2, start2,timetot2
call cpu_time (start2)!initial time
do ii=1,nn
do jj=1,nn
do kk=1,nn
mris2(jj,ii)=mris2(jj,ii)+m1(jj,kk)*m2(kk,ii)
!SCALAR PRODUCT BTW
!jj-th ROW OF m1 AND ii-th COLUMN OF m2
!IT GOES IN THE ENTRY (j,i) OF mris2
end do
end do
end do
call cpu_time (finish2)!finish time
timetot2=finish2-start2
end subroutine colcolmatmul

subroutine matmulintrinsic(nn,m1,m2,mris3,timetot3)
!third method: intrinsic matmul
implicit none
integer*2 ii, jj, kk, nn
integer*2, dimension(nn,nn) :: m1
integer*2, dimension(nn,nn) :: m2
integer*2, dimension(nn,nn) :: mris3
real :: finish3, start3,timetot3

call cpu_time (start3)!initial time
mris3= matmul(m1,m2)
!matmul DOES A MATRIX-MATRIX MULTIPLICATION m1*m2
call cpu_time (finish3)!finish time
timetot3= finish3-start3
end subroutine matmulintrinsic

end module matmultiplications
```

```

!MODULE DEBUGGING
module debugging

interface check
module procedure checkdim,checkmat
end interface

contains

subroutine checkdim(nn,nncheck,debug) !checking
    dimensions nn
implicit none
integer*2 :: nn, nncheck
logical :: debug

if(debug.eqv..TRUE.) then

if(nn>10000) then !is the dim too large?
print*, "WARNING:_too_large_dimension:"
print*, nn
else if(nn<1) then !is the dim minor than 1?
print*, "WARNING:_dimension_minor_than_1"
else if((nn-nncheck)>0.5) then !is the dim wrong?
print*, "the_dimension_is_wrong,_it_should_be:"
$                                     , nncheck
print*, "but_is", nn
print*, "_"
else
print*, "okay:_right_dimensions_matrices",nn
end if
end if
end subroutine checkdim

subroutine checkmat(nn,m,mcheck,debug)
implicit none
!checking if the input matrixes are equal
integer*2, dimension(nn,nn) :: m, mcheck
integer*2 :: tt,ss,nn
integer*4 accum
logical debug

if(debug.eqv..TRUE.) then
accum=0
do tt=1,nn
do ss=1,nn
if(m(tt,ss)/=mcheck(tt,ss)) then

```



```
        accum=accum+1
    end if
end do
end do

if(accum>0) then
print*, "The_two_matrices_have"
$, accum, "different_entries"
else
print*, "Same_matrices"
end if

end if
end subroutine checkmat

end module debugging

program test_performance_mulmat
!This program implements matrix-matrix multiplication
!in three different method: column by column, row by
    row
!and by intrinsic function. For each method the
    CPU_TIME
!is computed. This is done for squared matrix, from
!100x100 to 1000x1000, with a step=100.
!The results are printed on three different files,
    according
!to the used method.

!DECLARATION VARIABLES

use matmultiplications
use debugging
implicit none

integer*2, dimension(:,,:), allocatable :: m1
integer*2, dimension(:,,:), allocatable :: m2
integer*2, dimension(:,,:), allocatable :: mcheck1,
    mcheck2
integer*2, dimension(:,,:), allocatable :: mris1,
    mris2, mris3
integer*2 ii, jj, nn, nncheck
real :: ttime
real :: timetot1, timetot2, timetot3
```

```
real :: t1,t2, accumtime, t3,t4
logical :: debug
character*1 :: choice
integer :: my_stat
character (256) :: my_msg
call cpu_time(t1)!INITIAL TIME ALL PROGRAM

accumtime=0.0
choice="n" !Initialization of choice variable
!if it is "X", it asks the debug
!otherwise the programmer can choose between
!doing the debug or not writing "y" or "n"

if(choice=="y") then
debug=.TRUE.
else if(choice=="n") then
debug=.FALSE.
else
print*, "Not_understood"
stop
end if

!OPENING THE 3 "Result" FILE

!First method:row by row
open(unit = 90, file = "Results-1",
$      status = "unknown", access="append",
$  iostat=my_stat, iomsg=my_msg)

if(my_stat /= 0) then
print*, 'Open_Results-1_failed_with_stat_='
$      , my_stat, '_msg_='//trim(
    my_msg)
end if

!Second method:col by col
open(unit =80, file = "Results-2",
$      status = "unknown", access="append",
$  iostat=my_stat, iomsg=my_msg)

if(my_stat /= 0) then
print*, 'Open_Results-2_ failed_with_stat_='
```

```
$                                , my_stat, '_msg_=' //trim(
    my_msg)
end if

!Third method:intrinsic matmul
open(unit = 70, file = "Results-3",
$      status = "unknown", access="append",
$      iostat=my_stat, iomsg=my_msg)

if(my_stat /= 0) then
print*, 'Open-Results-3_failed_with_stat_='
$                                , my_stat, '_msg_=' //trim(
    my_msg)
end if

!HERE THE MATRIX DIMENSION IS TAKEN AS IMPUT FROM
!FILE "MatDimension.txt"

open(unit = 30, file = "MatDimension.txt",
$      status = "unknown",
$      iostat=my_stat, iomsg=my_msg)

if(my_stat /= 0) then
print*, 'Open-MatDimension-_failed_with_stat_='
$                                , my_stat, '_msg_=' //trim(
    my_msg)
end if

read(30,*) nn
nncheck=nn

!ALLOCATION MATRICES
!DEBUG
if(debug.eqv..TRUE.) then!DEBUG
print*, "_"
print*, "MATRICES_OF_ORDER_", nn
print*, "_"
end if

if(debug.eqv..TRUE.) then!checkdim
call check(nn,nncheck,debug)
end if
```

```

allocate(m1(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating m1
if(my_stat /= 0) then
print*, 'Failed_to_allocate_m1_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if
allocate(m2(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating m2
if(my_stat /= 0) then
print*, 'Failed_to_allocate_m2_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if
allocate(mris1(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating mris1
if(my_stat /= 0) then
print*, 'Failed_to_allocate_mris1_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if

allocate(mris2(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating mris2
if(my_stat /= 0) then
print*, 'Failed_to_allocate_mris2_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if
allocate(mris3(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating mris3
if(my_stat /= 0) then
print*, 'Failed_to_allocate_mris3_with_stat_='
$           , my_stat, '_and_msg_='//trim(my_msg
    )
end if

allocate(mcheck1(nn,nn), stat=my_stat, errmsg=my_msg)
!allocating mcheck1
if(my_stat /= 0) then
print*, 'Failed_to_allocate_mcheck1_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if
allocate(mcheck2(nn,nn), stat=my_stat, errmsg=my_msg)

```

```
!allocating mcheck2
if(my_stat /= 0) then
print*, 'Failed_to_allocate_mcheck2_with_stat_='
$           , my_stat, '_and_msg_='//trim(
    my_msg)
end if

!INITIALIZATION m1: in the entry (i,j)
!the value i+j is assigned
do ii=1,nn
do jj=1,nn
m1(ii,jj)=ii+jj
mcheck1(ii,jj)= m1(ii,jj)
end do
end do
!DEBUG
call cpu_time(t3)
if(debug.eqv..TRUE.) then
print*, "_"
print*, "INITIALIZATION_m1"
end if
if(debug.eqv..TRUE.) then!checkdim
call check(nn,nncheck,debug)
end if

!INITIALIZATION m2, in the entry (i,j)
!the value i*j is assigned
do ii=1,nn
do jj=1,nn
m2(ii,jj)=ii*jj
mcheck2(ii,jj)= m2(ii,jj)
end do
end do

!DEBUG
if(debug.eqv..TRUE.) then
print*, "_"
print*, "INITIALIZATION_m2"
end if
if(debug.eqv..TRUE.) then !checkdim
call check(nn,nncheck,debug)
end if
```

```

if(debug.eqv..TRUE.) then !check matrices
print*, "Imput_matrix_1"
call check(nn,m1,mcheck1,debug)
print*, "Imput_matrix_2"
call check(nn,m2,mcheck2,debug)
end if

!FIRST METHOD: row by row in resulting matrix

call rowrowmatmul(nn,m1,m2,mris1,timetot1)

write(90,"(I4,4X,E16.9)",iostat=my_stat, iomsg=my_msg
)
$
nn, timetot1

if(my_stat /= 0) then
print*, 'Write-Results-1- failed_with_stat_='
$, my_stat, '_msg_='//trim(my_msg)
end if
!printing on file the cpu_time

!DEBUG
if(debug.eqv..TRUE.) then
print*, "_"
print*, "FIRST_METHOD:_row_by_row"
end if
if(debug.eqv..TRUE.) then!checkdim
call check(nn,nncheck,debug)
end if

if(debug.eqv..TRUE.) then!checkmat
print*, "Imput_matrix_1"
call check(nn,m1,mcheck1,debug)
print*, "Imput_matrix_2"
call check(nn,m2,mcheck2,debug)
end if

!SECOND METHOD: column by column in resulting matrix

call colcolmatmul(nn,m1,m2,mris2,timetot2)

```

```

write(80,"(I4,4X,E16.9)",iostat=my_stat, iomsg=my_msg
)
$          nn, timetot2

if(my_stat /= 0) then
print*, 'Write_Results-2-_failed_with_stat_='
$          , my_stat, '_msg_='//trim(my_msg)
end if
!printing on file the cpu_time

!DEBUG

if(debug.eqv..TRUE.) then
print*, "_"
print*, "SECOND_METHOD:_col_by_col"
end if
if(debug.eqv..TRUE.) then!checkdim
call check(nn,nncheck,debug)
end if
if(debug.eqv..TRUE.) then!checkmat
print*, "Imput_matrix_1"
call check(nn,m1,mcheck1,debug)
print*, "Imput_matrix_2"
call check(nn,m2,mcheck2,debug)
end if

!THIRD METHOD: intrinsic function matmul

call matmulintrinsic(nn,m1,m2,mris3,timetot3)

write(70,"(I4,4X,E16.9)",iostat=my_stat, iomsg=my_msg
)
$          nn, timetot3

if(my_stat /= 0) then
print*, 'Write_Results-3-_failed_with_stat_='
$          , my_stat, '_msg_='//trim(my_msg)
end if
!printing on file the cpu_time

!DEBUG

```

```

if(debug.eqv..TRUE.) then
print*, " "
print*, "THIRD_METHOD:_intrinsic_matmul"
end if
if(debug.eqv..TRUE.) then!checkdim
call check(nn,nncheck,debug)
end if
if(debug.eqv..TRUE.) then!checkmat
print*, "Imput_matrix_1"
call check(nn,m1,mcheck1,debug)
print*, "Imput_matrix_2"
call check(nn,m2,mcheck2,debug)
end if

!the resulting matrices are different?
!DEBUG
if(debug.eqv..TRUE.) then
print*, " "
print*, "CHECKING_RESULTING_MATRICES"
end if
if(debug.eqv..TRUE.) then!checkdim
print*, "Checking_first_and_second_method"
call check(nn,mris1,mris2,debug)
print*, "Checking_third_and_second_method"
call check(nn,mris2,mris3,debug)
print*, "Checking_first_and_third_method"
call check(nn,mris1,mris3,debug)
end if

!DEALLOCATION matrices:

deallocate(m1, stat=my_stat, errmsg=my_msg)
!deallocating m1
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_m1_with_stat=_ '
$                                , my_stat, ' _and_msg=_ '//
    trim(my_msg)
end if

deallocate(m2, stat=my_stat, errmsg=my_msg)
!deallocating m2
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_m2_with_stat=_ '

```



```
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

deallocate(mris1, stat=my_stat, errmsg=my_msg)
!deallocating mris1
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_mris1_with_stat_='
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

deallocate(mris2, stat=my_stat, errmsg=my_msg)
!deallocating mris2
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_mris2_with_stat_='
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

deallocate(mris3, stat=my_stat, errmsg=my_msg)
!deallocating mris3
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_mris3_with_stat_='
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

deallocate(mcheck1, stat=my_stat, errmsg=my_msg)
!deallocating mcheck1
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_mcheck1_with_stat_='
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

deallocate(mcheck2, stat=my_stat, errmsg=my_msg)
!deallocating mcheck2
if(my_stat /= 0) then
print*, 'Failed_to_deallocate_mcheck2_with_stat_='
$                                , my_stat, '_and_msg_=' //
    trim(my_msg)
end if

stop
```

```
end program test_performance_mulmat
```